

Information Retrieval Project

Elise Kuylen Robin Verachtert

1 Introduction

For this project, we implemented a *content-based* recommendation system for books. Based on which books a user has liked before, we tried to recommend new books, that were similar in content and would (hopefully) also interest the user.

2 Algorithms implemented

To create our recommendation system, we first had to index a set of documents - the summaries and reviews of the books in our database. Next, we constructed tf-idf vectors for these books, which we could then compare to determine the similarity of different books. We also implemented dimensionality reduction.

2.1 Indexing a collection of documents

2.2 Vector space model

2.3 Dimensionality reduction

To perform a dimensionality reduction we had to use Singular Value Decomposition. For this we consulted the chapter in the book describing the dimensionality reduction. /

3 Design choices

3.1 Dataset collection

To get a sufficiently large dataset of books and user reviews of those books, we scraped the Goodreads database. Goodreads is a social platform that allows users to rate and review the books they have read.

We saved the information we got from Goodreads in JSON format, since this is easy to parse. We collected some different datasets, of different sizes, so we had a good variety to test the project on. There is a small dataset, only containing books of 2 users, this is the smallest set, which we used to quickly test newly implemented parts. Then there is a bookset of about 2000 books and one of about 10000 books. The last one we mostly used to test the speed of the indexer, because on our normal computers the similarity computation took too long.

We also have an additional dataset containing reviews of clothing, shoes, and jewelry from Amazon ¹. This is a larger dataset containing 250000 reviews.

¹ <http://jmcauley.ucsd.edu/data/amazon/>

This dataset is exclusively used for testing the speed of the indexer.

3.2 Lucene for indexing

To index the summaries and reviews of the books that we had scraped from Goodreads, we used the Lucene library.

3.3 Creating tf-idf vectors

To store the TF-IDF vectors of our books, we create the class `TFIDFBookVector`.

3.4 Recommendation

To make recommendations we used two different strategies:

- **Based on one book:** To make recommendations based on one book in the user's profile, we treated the text of the summary and the reviews for this book as a query. Next, we constructed a TF-IDF vector for this query, and compared it to the TF-IDF vectors in our matrix, using the cosine similarity measure. Once this was done, we sorted the books according to their similarity with the TF-IDF vector for the user's book and returned the first 10 results.
- **Based on all books in the user's profile**

3.5 Dimensionality reduction

We decided to not implement this ourselves, because there are enough libraries supporting it.

The SVD library we used is UJMP ². This also provided a sparse matrix implementation which is relevant to the total memory consumed. The implementation is very straightforward. You start with a matrix $M \times N$ and perform the SVD which provides 3 matrixes. U ($M \times M$), Σ ($M \times N$) and V ($N \times N$). Then you reduce the complexity by keeping the k first columns of U and V and making Σ a $k \times k$ matrix with the k highest singular values on the diagonal. From this you can compute the term matrix again (in the reduced space) by doing:

$$M_k = U_k * \Sigma_k * V_k^T$$

When we want to match a query we have to also map this query vector to the reduced dimension space.

$$q_k = \Sigma_k^{-1} * U_k^t * q$$

This gives a vector of the same size as q , so we can use it the same as before to query against the term Matrix.

² <https://ujmp.org/>

4 Evaluation results

4.1 Indexing

To evaluate the quality of our indexing process, we timed how long it took for different sizes of datasets. We checked the time consumption for regular indexing as well as for indexing followed by dimensionality reduction.

Dataset	Time Consumption
Clothing,Shoes,and Jewelry data set from Amazon (150MB) ³	45388 ms
10000 books dataset	9332 ms
180 books dataset	2800 ms

4.1.1 Regular

4.1.2 With dimensionality reduction

4.2 Recommendation

To evaluate our recommendation process, we looked at the precision of our results.⁴ To check the precision of the results, we looked at some users with a sizeable amount of books in their profile. 80% of these we used as training set, to base our recommendations on. The other 20% we used for evaluation. Since we are working with ranked recommendation (on the basis of how similar texts are), we used precision @ x to evaluate. Since giving a top-10 is an often-used strategy in recommendation, we will evaluate the precision @ 10 of our results.

4.2.1 Regular

4.2.2 With dimensionality reduction

5 External sources

6 Code

The code created in this project can be found at https://github.com/verachtertr/IR_project

⁴ Recall is less relevant here, and it is also nigh impossible to check: we would have to decide manually, for each book in our dataset, whether a book is relevant for a particular user or not.